

# 4. Communication (Part I)

Engineering Web and Data-intensive Systems

**Dr. Volker Riediger - Winter Term 2022/23**

# Communication (Part I)

- Network basics
- HTTP
- Sessions
- Scaling and Load Balancing



Image: colourbox.de

# Chapter 4.

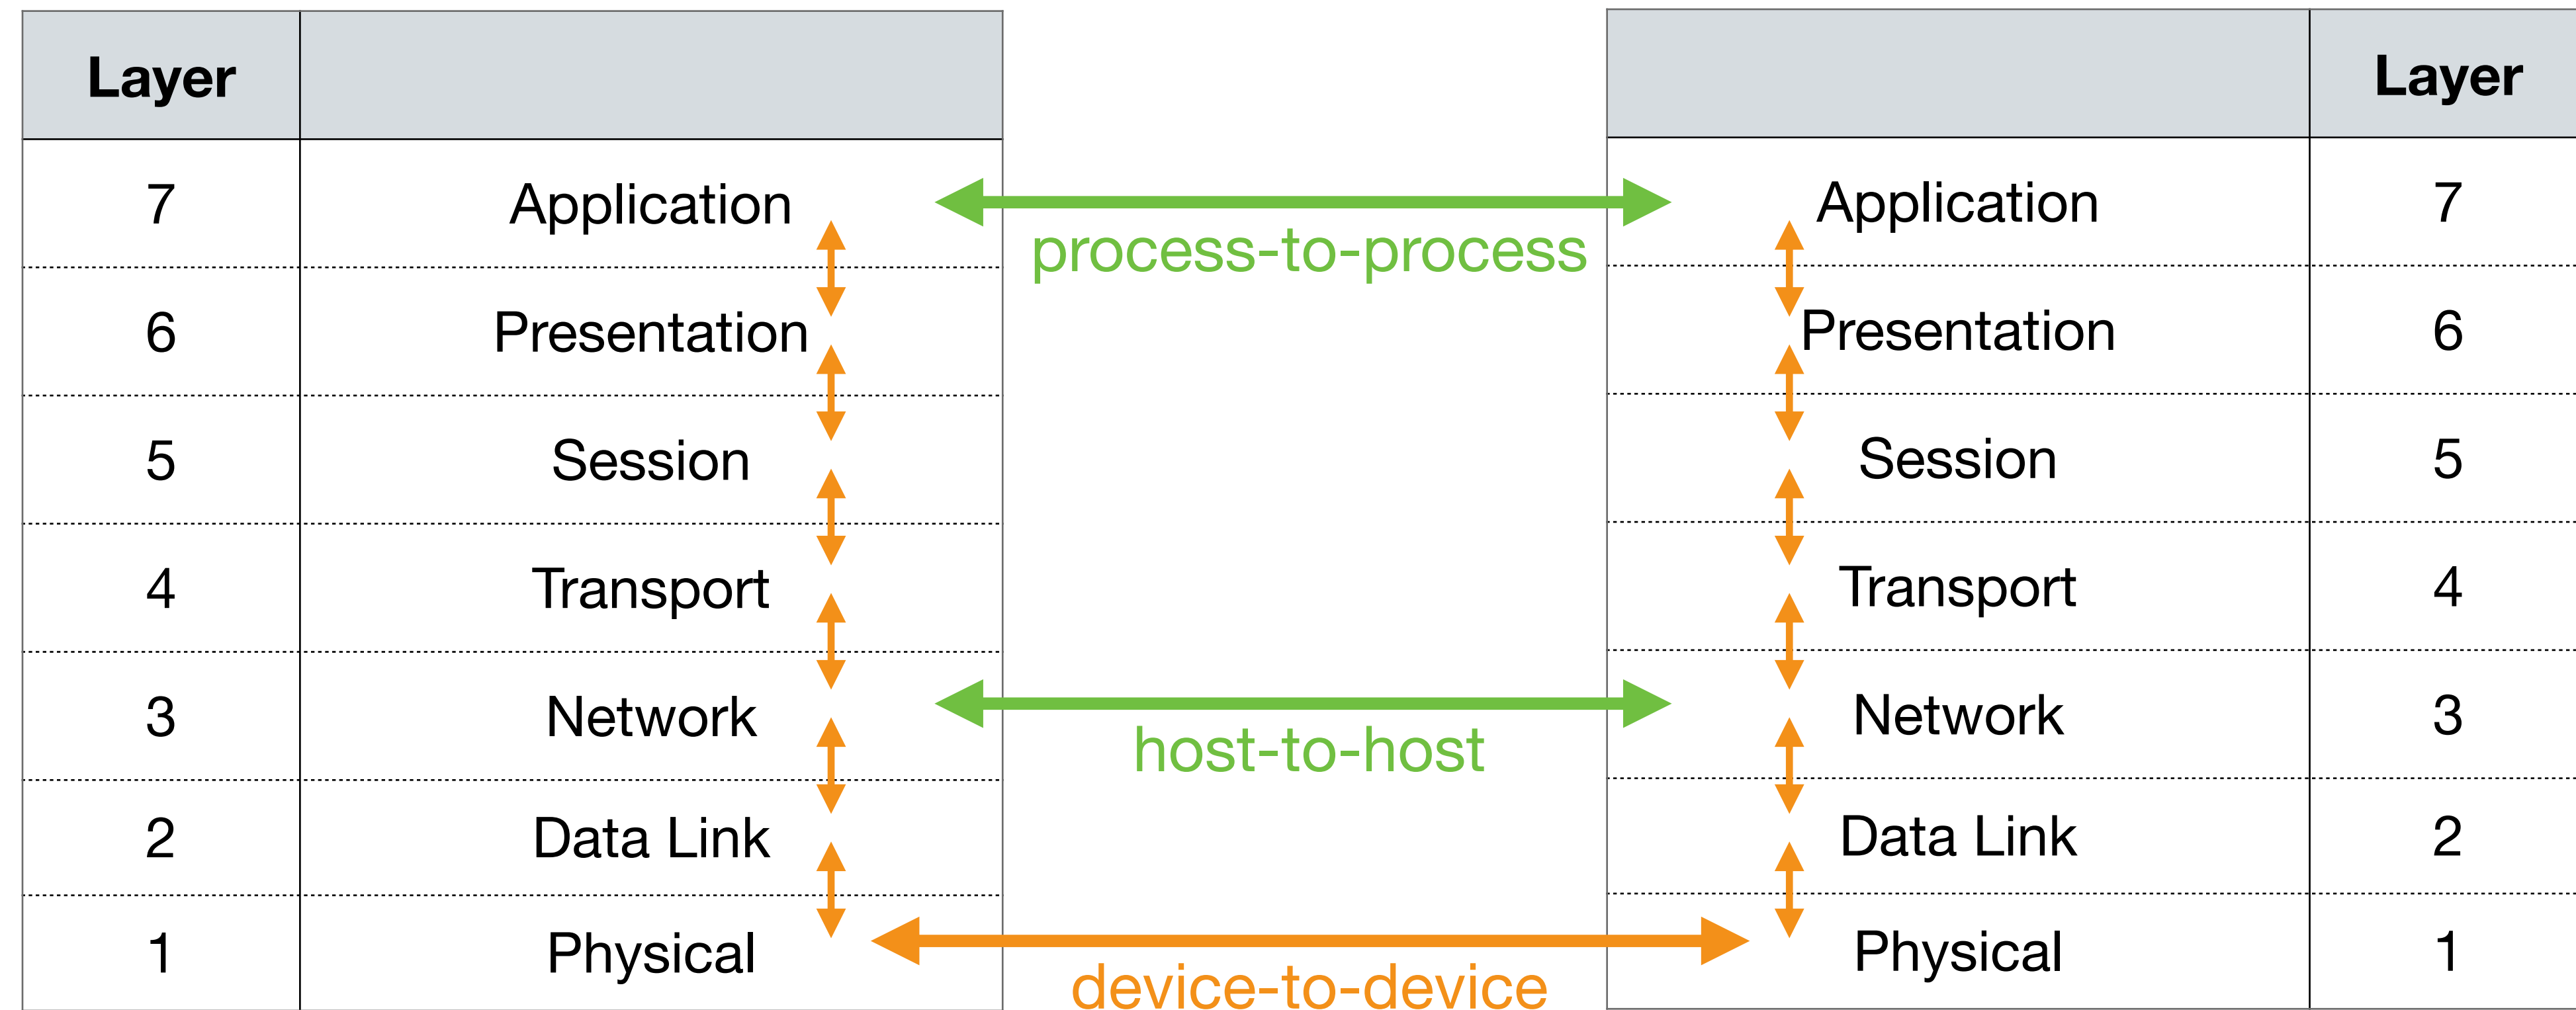
# Communication



# 4.1 Network Communication

# Network Communication

- To enable communication between network nodes, models, architecture, rules, logical, and physical properties have to be agreed upon by the communicating parties
  - Reference architecture defined in the **Open Systems Interconnection** (OSI) project conducted by the International Standard Organization (ISO)
  - Many **implementations** emerged for different communication scenarios, e.g.
  - TCP/IP - the internet protocol, transmission control protocol
  - IPX/SPX - Internetwork/Sequenced Packet Exchange
  - SNA - IBM Systems Network Architecture
  - UMTS - Universal Mobile Telecommunications System
  - Various **different physical transfer technologies** exist
- ➡ Need to hide the underlying complexity from the communicating parties

# OSI Reference Model



 logical communication  
 physical communication

# OSI Layers

OSI model				
Layer		Protocol data unit (PDU)	Function <sup>[6]</sup>	
Host layers	7	Application	Data	High-level APIs, including resource sharing, remote file access
	6	Presentation		Translation of data between a networking service and an application; including character encoding, data compression and encryption/decryption
	5	Session		Managing communication sessions, i.e., continuous exchange of information in the form of multiple back-and-forth transmissions between two nodes
	4	Transport	Segment, Datagram	Reliable transmission of data segments between points on a network, including segmentation, acknowledgement and multiplexing
Media layers	3	Network	Packet	Structuring and managing a multi-node network, including addressing, routing and traffic control
	2	Data link	Frame	Reliable transmission of data frames between two nodes connected by a physical layer
	1	Physical	Symbol	Transmission and reception of raw bit streams over a physical medium


from [\[https://en.wikipedia.org/wiki/OSI\\_model\]](https://en.wikipedia.org/wiki/OSI_model)



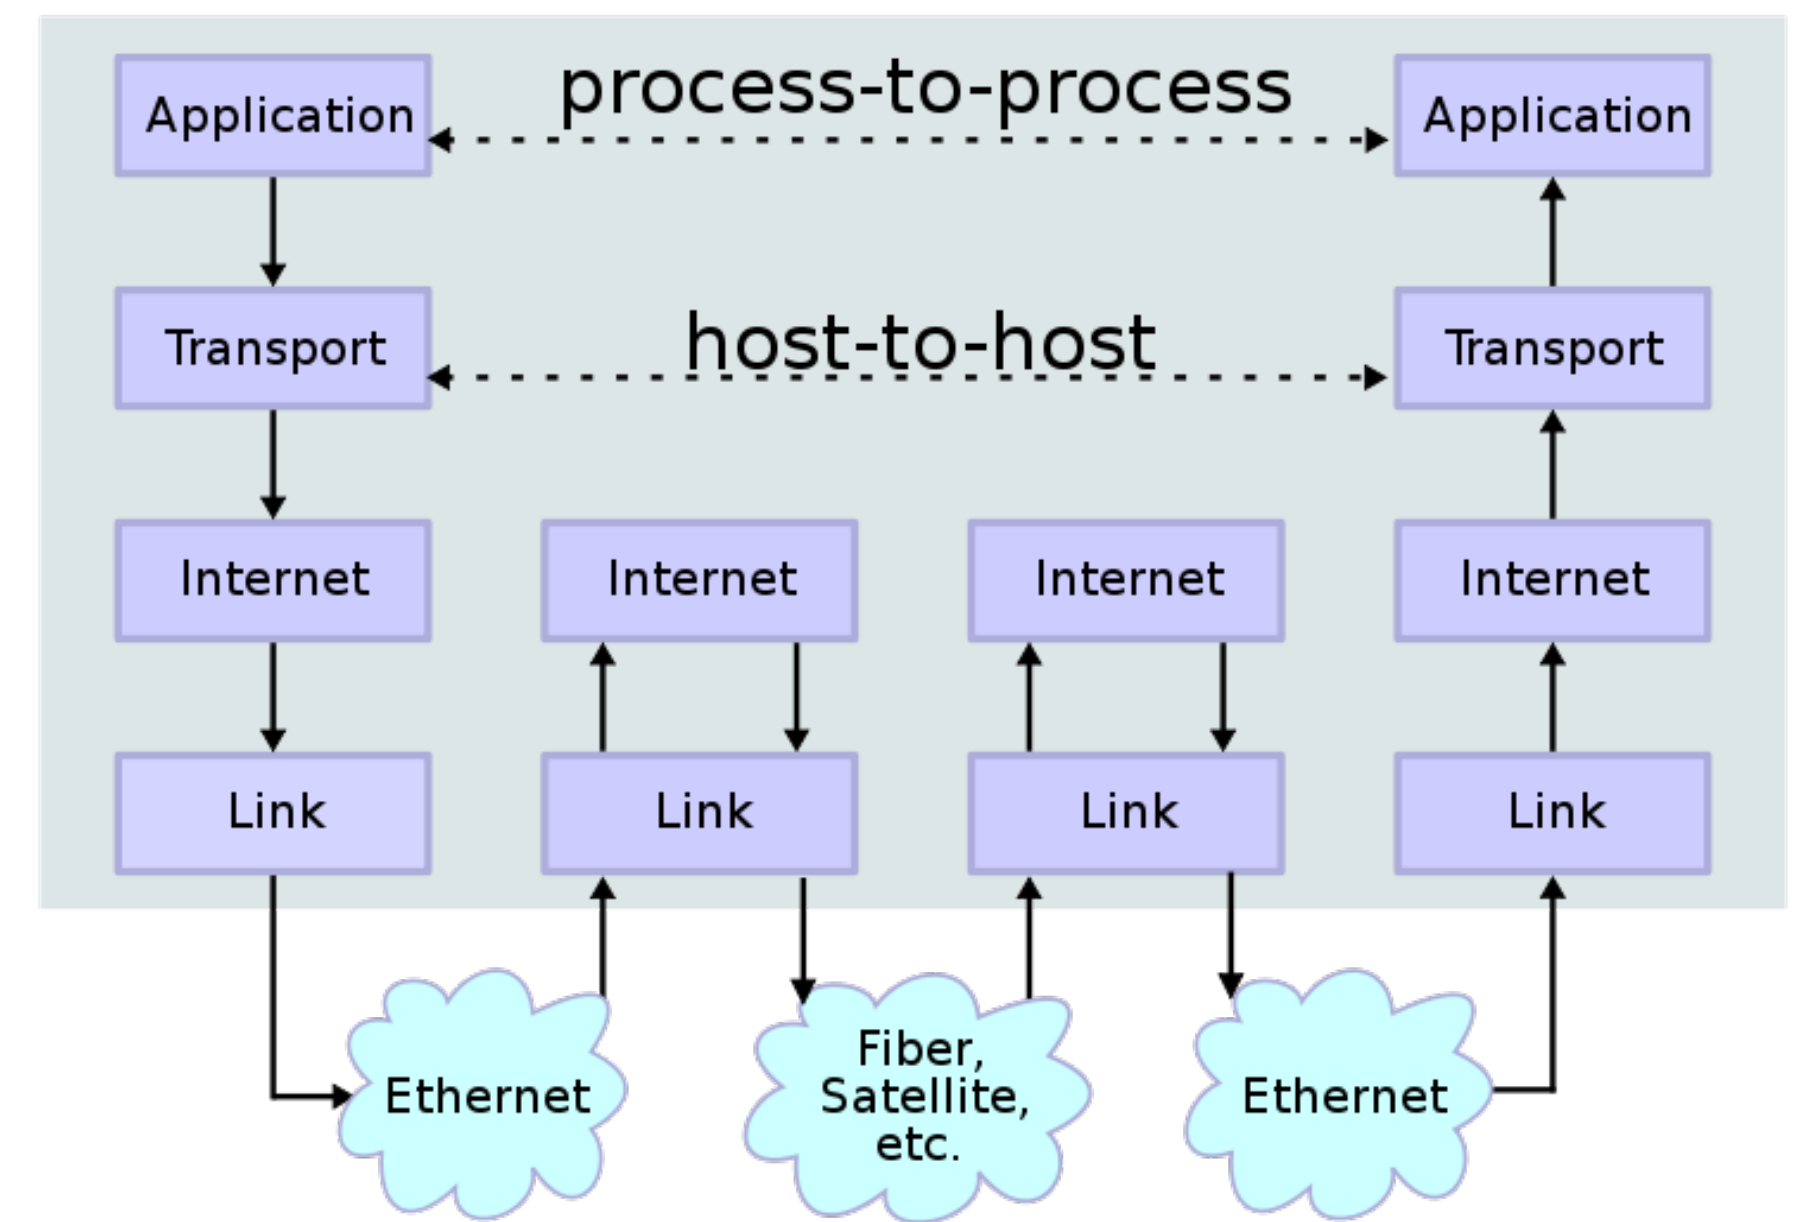
# TCP/IP

- Layered Network Protocol Scheme
- Layers relate to, but not exactly match, the OSI layers
- Offers 2 Protocols:
  - **UDP** User Datagram Protocol  
connectionless best effort transfer
  - **TCP** Connection-based  
bi-directional reliable transfer
- IP = Internet Protocol
  - Packet transfer between nodes
  - Address schemes
  - Routing services

## Network Topology



## Data Flow




[\[https://en.wikipedia.org/wiki/Internet\\_protocol\\_suite\]](https://en.wikipedia.org/wiki/Internet_protocol_suite)



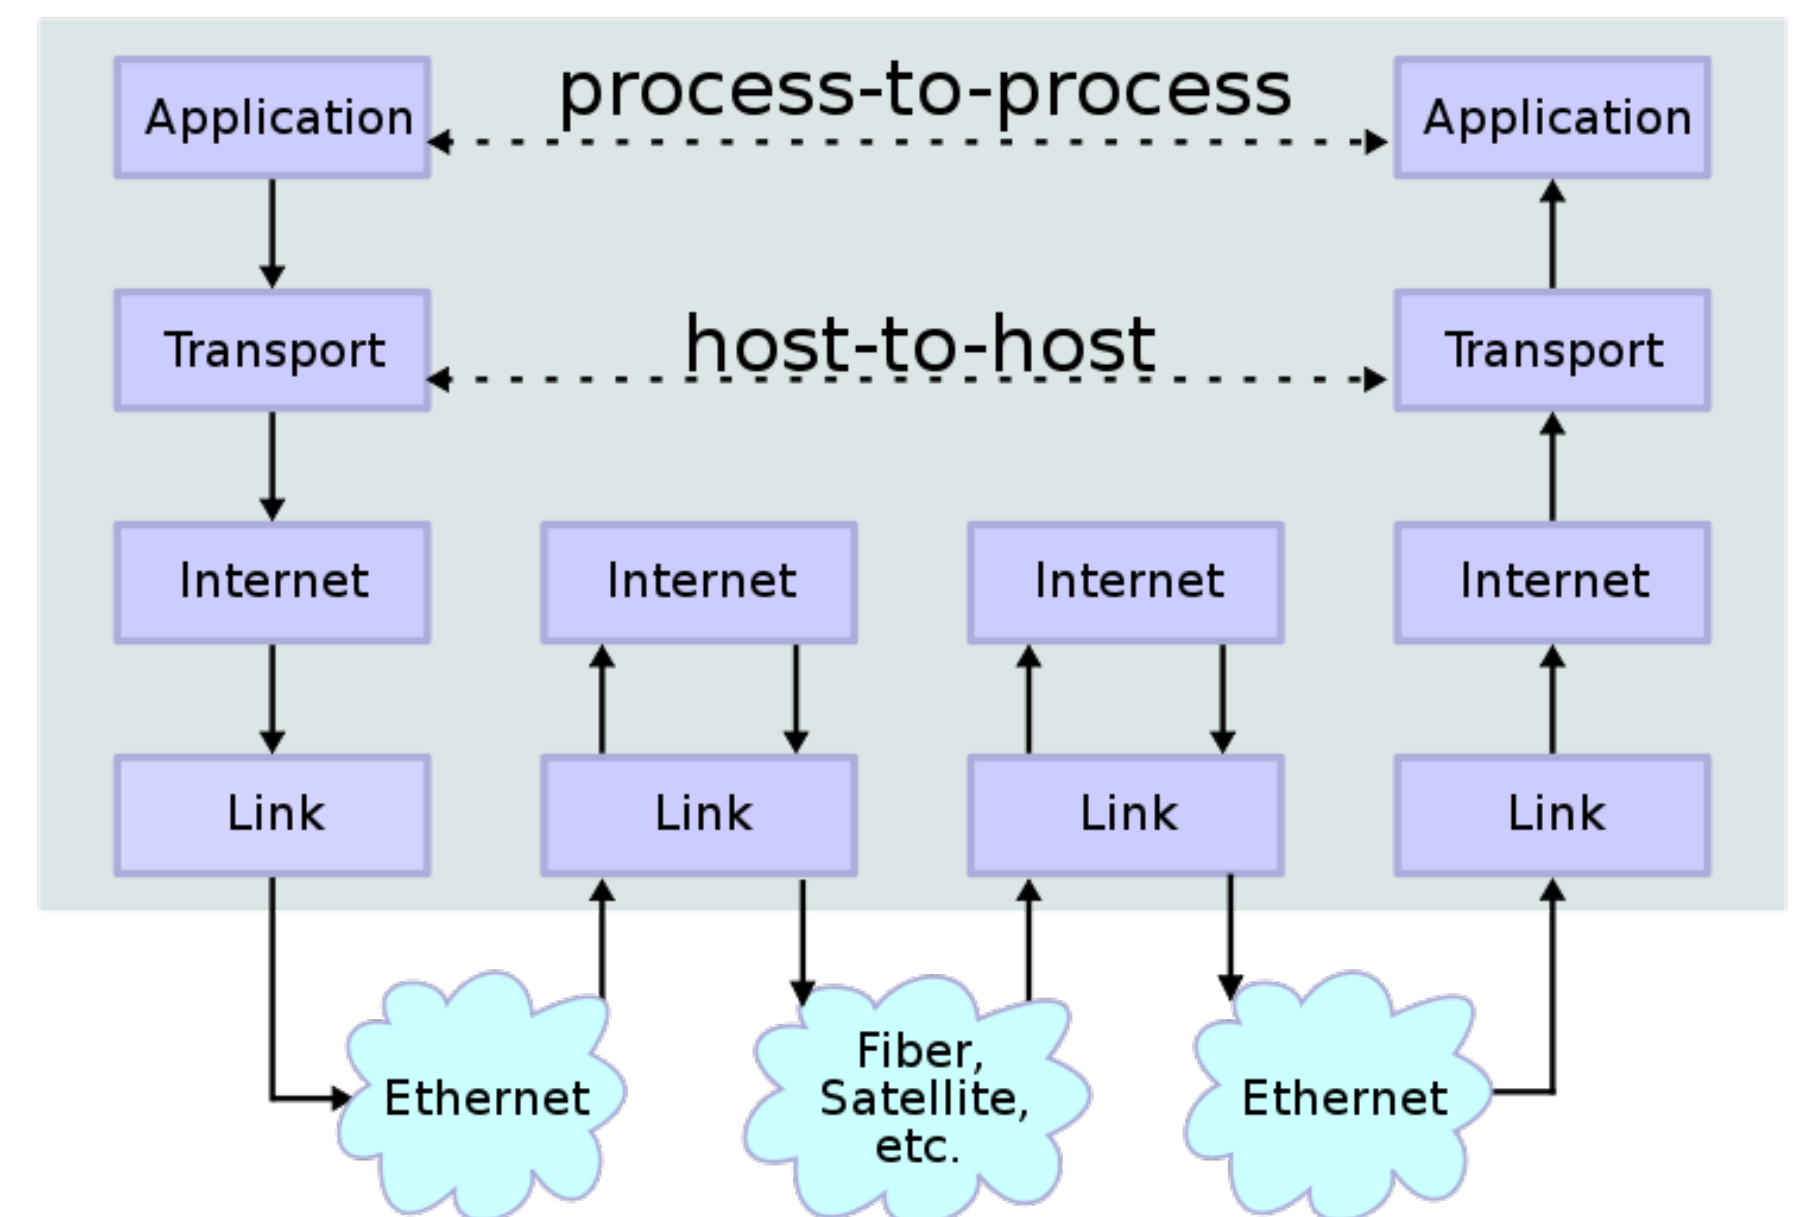
# TCP/IP

- On top of TCP, many **application protocols** are defined, e.g.
  - HTTP 80 / 443
  - SMTP 25 / 465
  - POP 110 / 965
  - IMAP 143 / 993
  - FTP 20,22 / 989,990
  - SSH 22
  - ...
- Protocols define **well-known ports** (part of UDP and TCP) to identify **endpoints**
- Registry managed by **IANA** (Internet Assigned Numbers Authority)

## Network Topology



## Data Flow



[https://en.wikipedia.org/wiki/Internet\\_protocol\\_suite](https://en.wikipedia.org/wiki/Internet_protocol_suite)

# 4.2 HTTP - Hyper Text Transfer Protocol

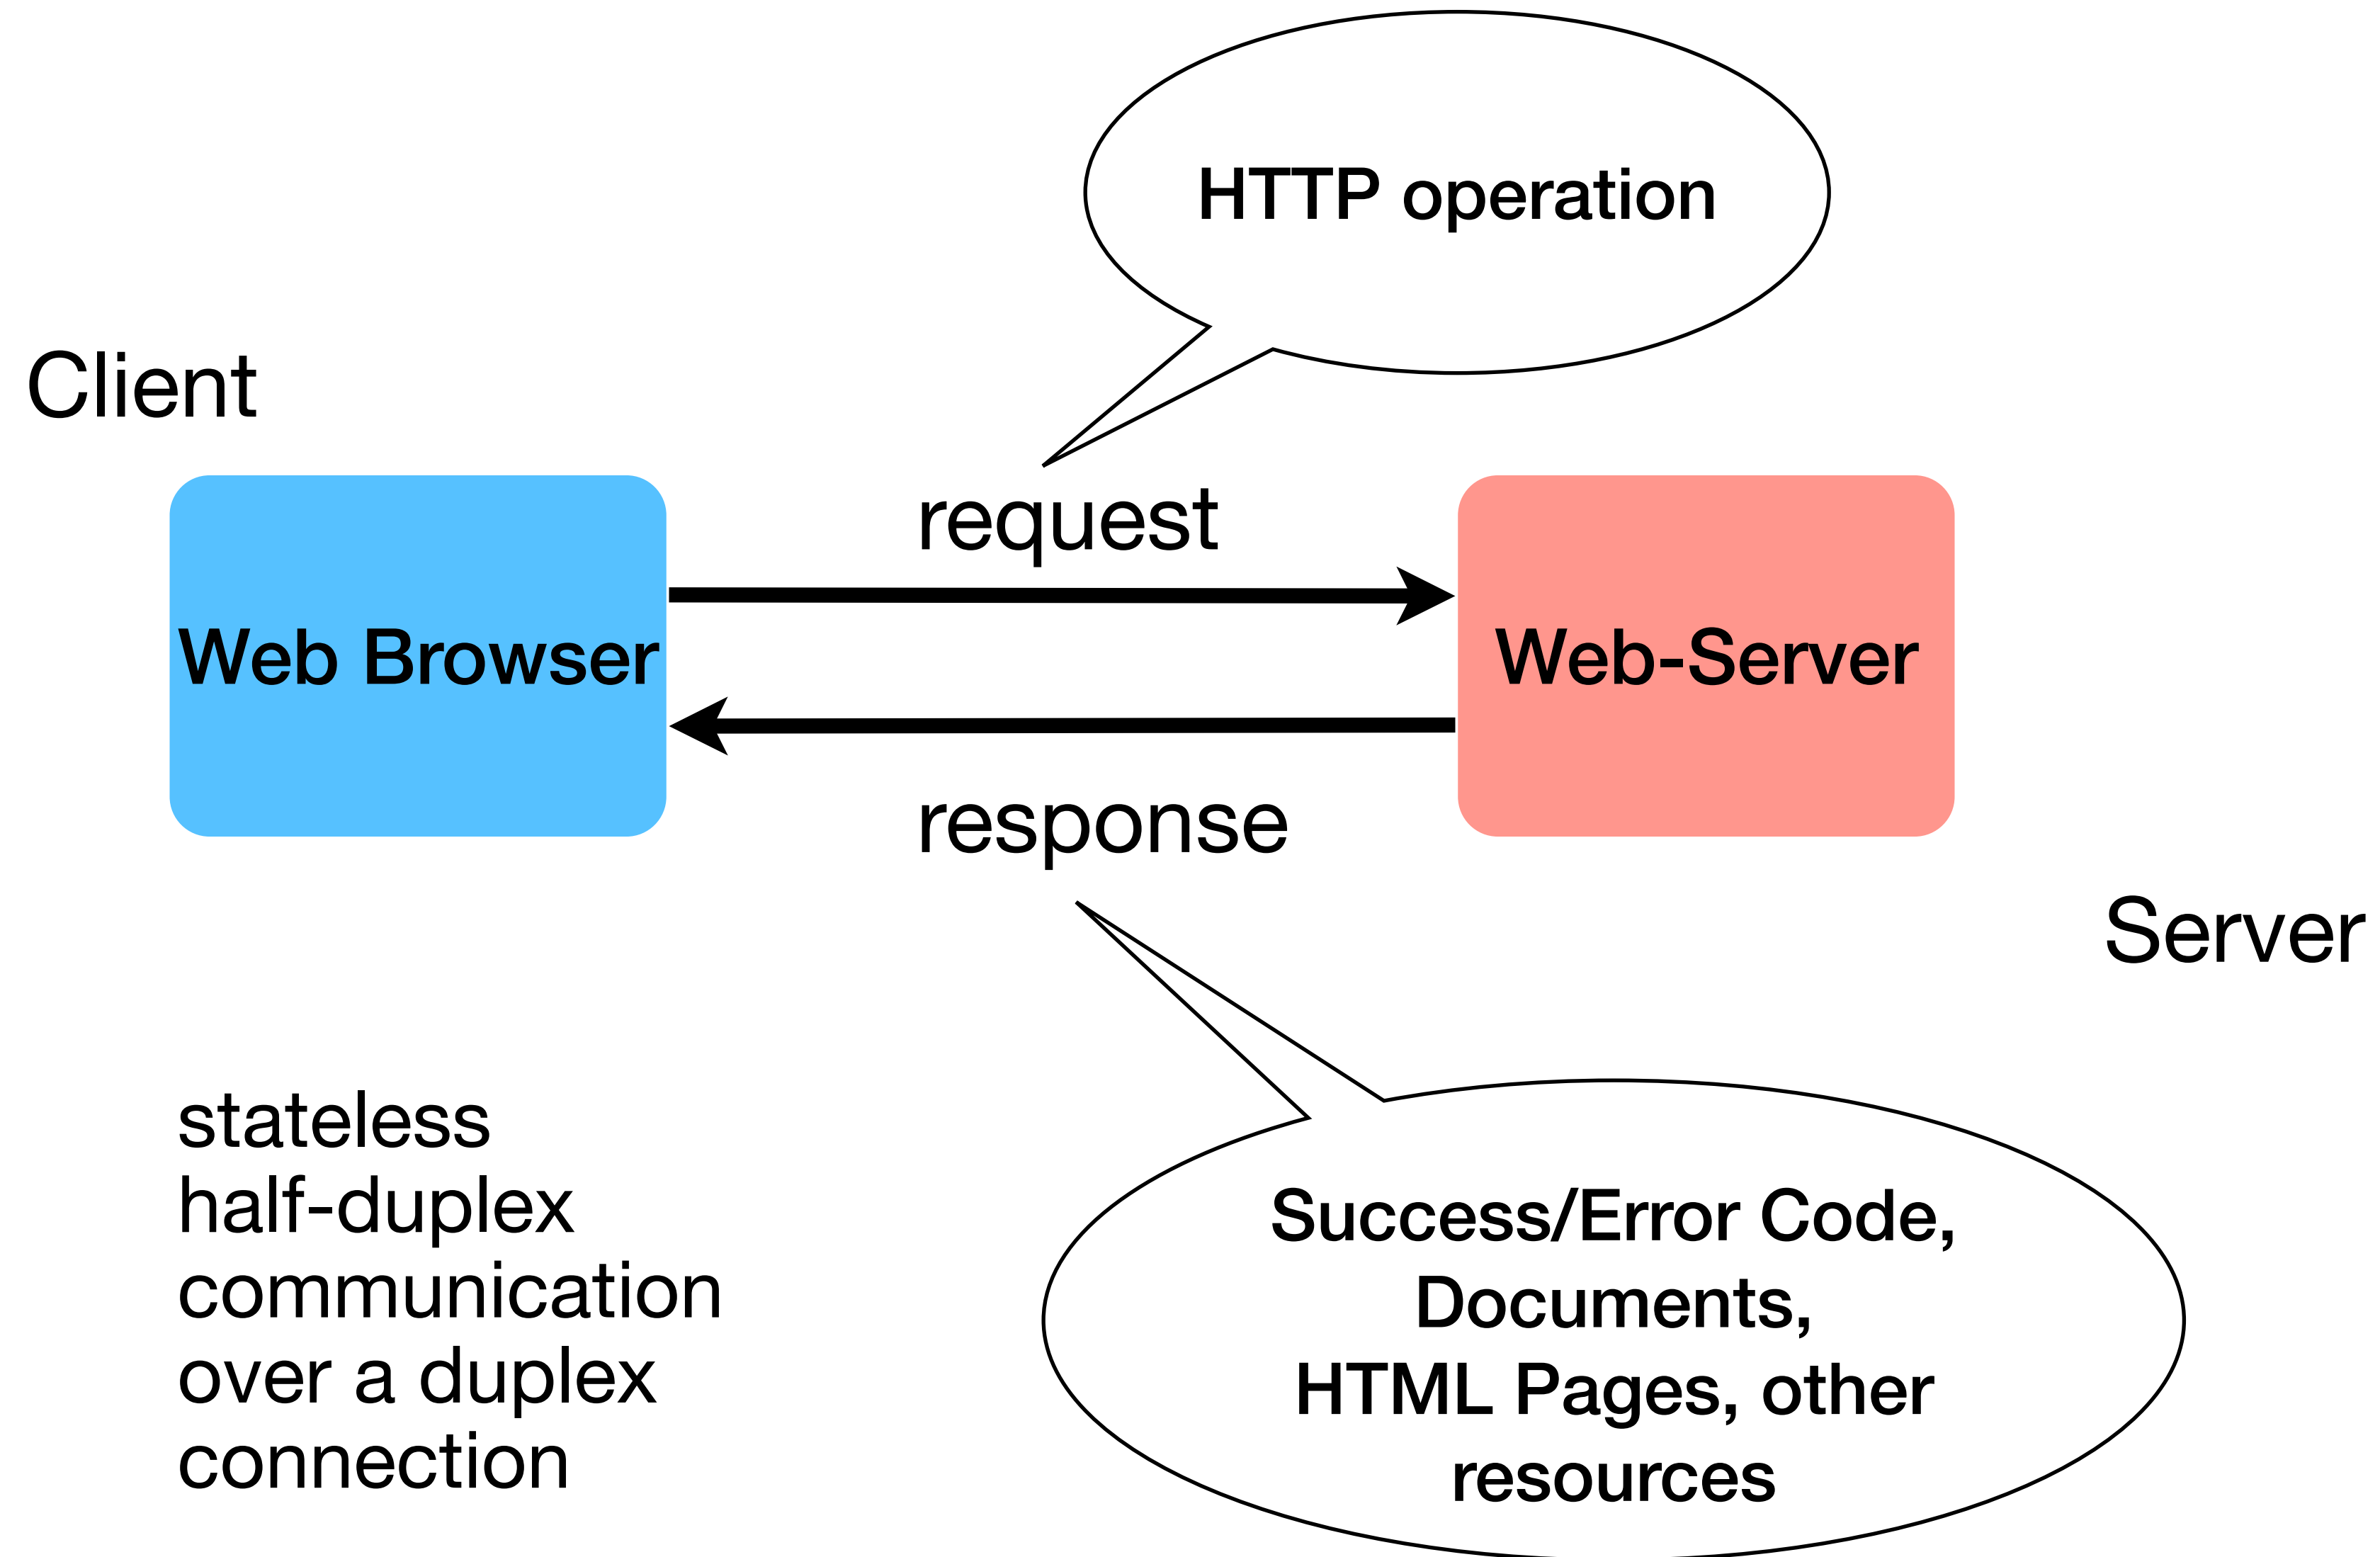
# Definitions

A **Web Server** is a program that delivers web pages using the Web using the **HTTP protocol**.

A **Web Browser** is a program that supports rendering of, and interaction with web pages written in HTML.

A **protocol** is a **formal description** of message formats and rules for **exchanging** those messages.

# Web Application: Client and Server



# HTTP

- Communication in the Web is done using the **Hypertext Transfer Protocol** (HTTP), also developed by Tim Berners-Lee.
- HTTP is a simple protocol on the **application layer** of the Internet, supporting the exchange of documents between the client and the server.
- HTTP is a coordinated result of IETF and W3C.

# HTTP in Web Applications

- transfer **static** content
  - text
  - stylesheets
  - pictures
  - ...
- transfer **dynamic** content depending on system state, e.g.
  - database content
  - time
  - user identity
  - ...
- provide **web services**
  - call business functions
  - modify data
  - ...

# HTTP Versions

Year	Version	Standard	Features
1991	0.9	-	early WWW prototype
1996	1.0	<a href="#">RFC 1945</a>	additional methods and header fields
1997	1.1	<a href="#">RFC 7231</a>	Multiple subsequent requests, CONNECT method
2015	2.0	<a href="#">RFC 7540</a>	request multiplexing i.e. multiple parallel requests, servers may push resources
2018	3.0	upcoming, draft status as of Dec 2019	support UDP Transport



# HTTP operations (methods)

- Method and its parameters are specified in the **request header**
  - **GET** loads a resource from a web server
  - **HEAD** loads meta information for a resource
  - **OPTIONS** list of allowed methods for a resource
  - **TRACE** echoes the command
  - **POST** resource specific processing of the payload, can also create
  - **PUT** replaces/creates a resource on the web server
  - **DELETE** removes a resource from the server
  - **CONNECT** (1.1) builds a tunnel when using proxies
  - **PATCH** (2.0) modifies (parts of) a resource
  - ...and many more, see also: [List of HTTP methods](#)
- these so-called **safe methods** are assumed to be **read-only and not to modify content**

# Form based Web Applications

## BROWSER

1. **Request** of a page (GET)

3. **Request** for further resources (GET)  
e.g. images, scripts, styles...

5. **Presentation** of the complete web page  
(rendering)

6. **Send** form data (POST)

repeated for each subsequent action

## WEB SERVER

2. **Delivery** of HTML document

4. **Delivery** of requested data  
steps 3. and 4. possibly repeated many times

7. **Processing** of form data, then delivery of  
response document  
(continue at step 2)

Please install `curl` for your operating system!  
Conduct your own experiments with curl.  
You'll have to use `curl` in the persistence assignments!

# Examples using `curl` and Browser Tools

# Single Page Web Applications

## WEB BROWSER

1. **Request** of a page (GET)
3. **Request** for further resources (GET)  
e.g. images, scripts, styles...
5. **Presentation** of the complete web page  
(rendering)

6. **Transmission** of a so-called  
„AJAX-Request“ (GET/POST/PUT/DELETE)

8. Processing of response,  
**modification** of the page,  
continue at step 6.

repeated for each subsequent action

## WEB SERVER

2. **Delivery** of HTML document

4. **Delivery** of requested data  
steps 3. and 4. possibly repeated many times

7. **Processing** of request,  
response as XML or JSON document, „pure data“

**Example:** RESToku

# HTTP Headers

- Request and response start with the **HTTP protocol header**
- **Request header:**
  - Method, URL path, and protocol version
  - Hostname (of the server)
  - optional information, e.g. content length, content encoding, cookies, accepted content type
  - a blank line
  - (after that, content may be transferred)

Example requesting `https://www.uni-koblenz.de`  
(The > and < marks are not part of the protocol but added by the cURL program)

```
> GET / HTTP/1.1
> Host: www.uni-koblenz.de
> User-Agent: curl/7.64.1
> Accept: */*
>
```

In this case, the server responds with a so-called redirect message, a web browser would then request the document indicated by the Location header. The response additionally contains human-readable HTML for browsers that don't automatically redirect.

```
< HTTP/1.1 301 Moved Permanently
< Date: Thu, 16 Jan 2020 12:06:13 GMT
< Server: Apache/2.4.18 (Ubuntu)
< Location: https://www.uni-koblenz-landau.de/de/koblenz/
< Content-Length: 338
< Content-Type: text/html; charset=iso-8859-1
<
<!DOCTYPE HTML PUBLIC "-//IETF//DTD HTML 2.0//EN">
<html><head>
<title>301 Moved Permanently</title>
<p>The document has moved <a href="https://www.uni-koblenz-landau.de/de/koblenz/">here</a>.</p>
<hr>
<address>Apache/2.4.18 (Ubuntu) Server at www.uni-koblenz.de
Port 443</address>
</body></html>
```



# HTTP Headers

- Request and response start with a **HTTP protocol header**
- **Response header:**
  - Protocol version, status code, status description
  - Timestamp
  - status code specific information
  - content length, content type, cookies, ...
  - a blank line
  - (after that, content may be transferred)

Example requesting `https://www.uni-koblenz.de`  
(The > and < marks are not part of the protocol but added by the cURL program)

```
> GET / HTTP/1.1
> Host: www.uni-koblenz.de
> User-Agent: curl/7.64.1
> Accept: */*
>
```

In this case, the server responds with a so-called redirect message, a web browser would then request the document indicated by the Location header. The response additionally contains human-readable HTML for browsers that don't automatically redirect.

```
< HTTP/1.1 301 Moved Permanently
< Date: Thu, 16 Jan 2020 12:06:13 GMT
< Server: Apache/2.4.18 (Ubuntu)
< Location: https://www.uni-koblenz-landau.de/de/koblenz/
< Content-Length: 338
< Content-Type: text/html; charset=iso-8859-1
<
<!DOCTYPE HTML PUBLIC "-//IETF//DTD HTML 2.0//EN">
<html><head>
<title>301 Moved Permanently</title>
<p>The document has moved <a href="https://www.uni-koblenz-landau.de/de/koblenz/">here</a>.</p>
<hr>
<address>Apache/2.4.18 (Ubuntu) Server at www.uni-koblenz.de
Port 443</address>
</body></html>
```



# Protocols “on top of HTTP”

- Private and commercial networks usually **block** most services from the internet, e.g. by placing **firewalls** (packet filters) between the external and the internal networks
- Many well-known services use **privileged port numbers**  $\leq 1023$
- Only a **few services** are exposed
  - Packets are **filtered** by firewalls based on source or destination IP addresses and/or protocol port numbers (and other criteria)
  - However, connections via HTTP(S) ports 80/443 are permitted in most networks

# Protocols “on top of HTTP”

- Hence, many application protocols were built “on top” of HTTP, thereby using the same port numbers and circumventing filter rules
- Commonly, at least the initiating communication is done via HTTP ports before using application specific ports
- The HTTP protocol has an [UPGRADE](#) method to switch to different protocols once the connection is established
- **WebDAV**  
HTTP extension for **D**istributed **A**uthoring and **V**ersioning ([RFC 4918](#))
- **WebSockets**  
HTTP extension (via [UPGRADE](#)) for bi-directional communication and push notifications ([RFC 6455](#))
- **Messaging**  
e.g. [Skype](#), [iMessage](#), [WhatsApp](#), and others
- ...

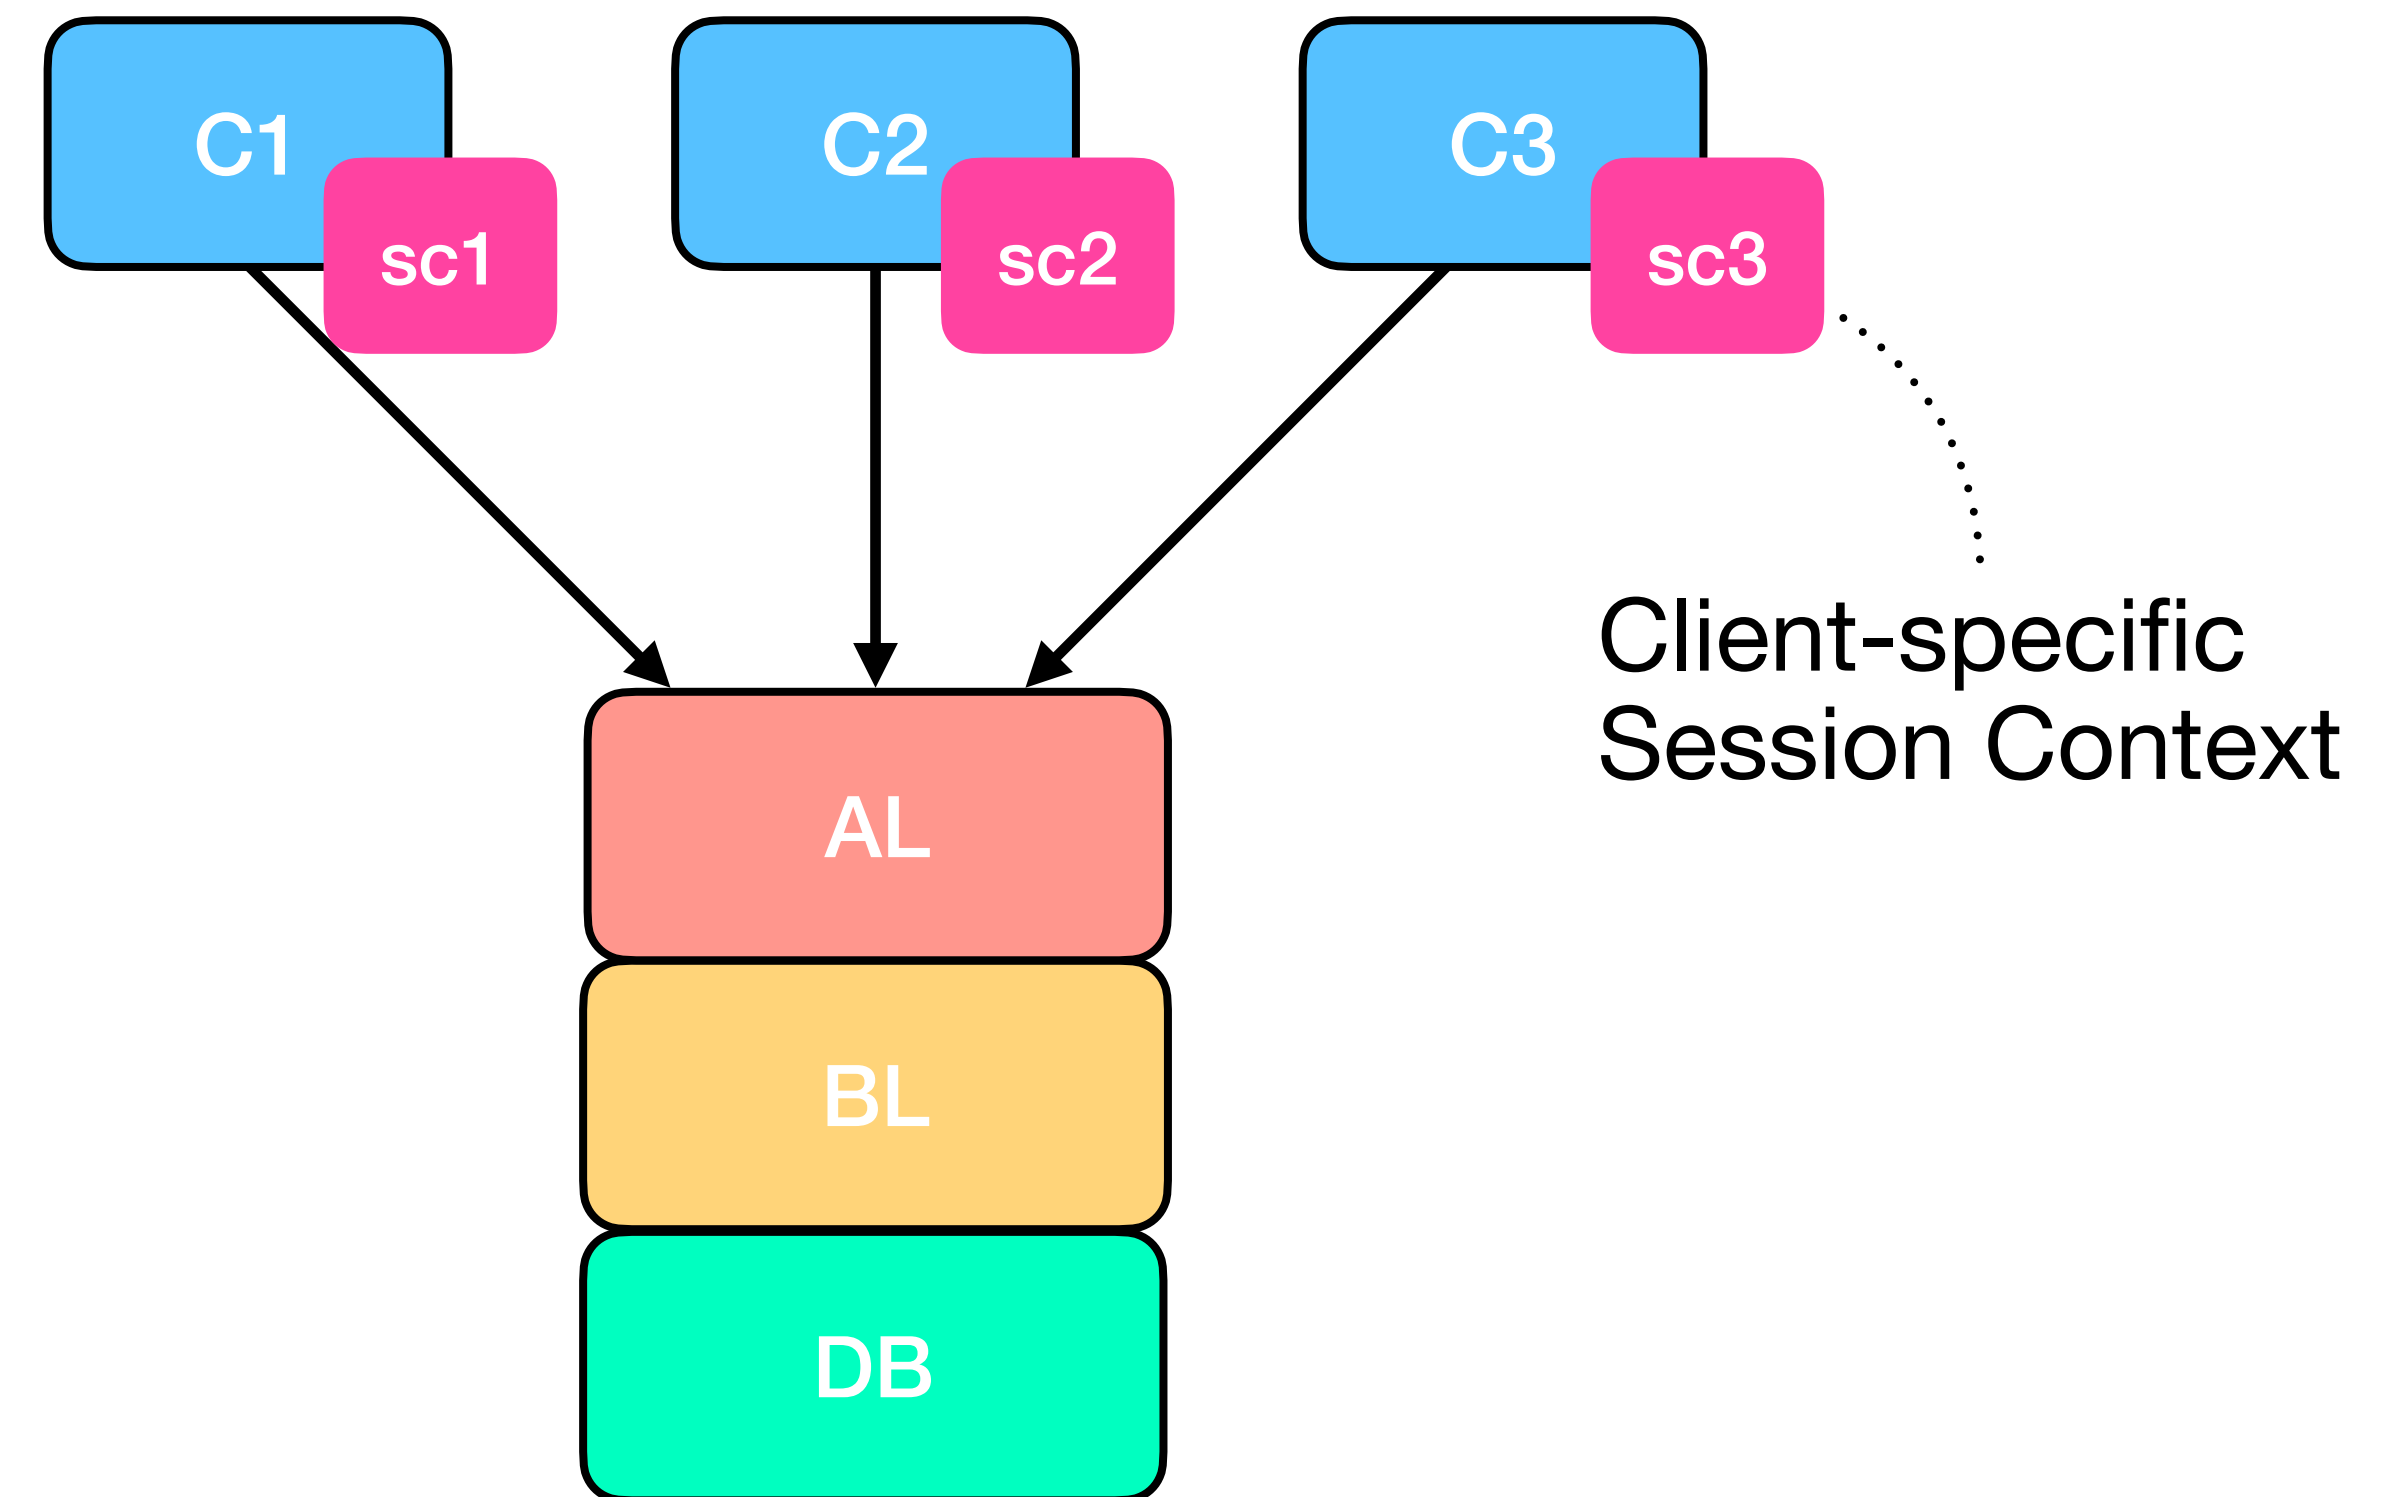
# 4.3 Sessions

# Sessions

- Problem: HTTP is **stateless** by definition
- Subsequent requests appear to a server as if they were **unrelated**, and HTTP servers *have to* handle requests as if they were unrelated.
- Web Applications require a **user session**, e.g. log-in, interact with the application, log-out
- Such a session has to “survive” many basic HTTP requests
- This requires that some **state data associated with the current client connection**, the so-called **session context**, has to be stored
- E.g. “current customer and associated shopping cart” in web shop systems
- Question: Where to store the session context?
  - Client side
  - Server side

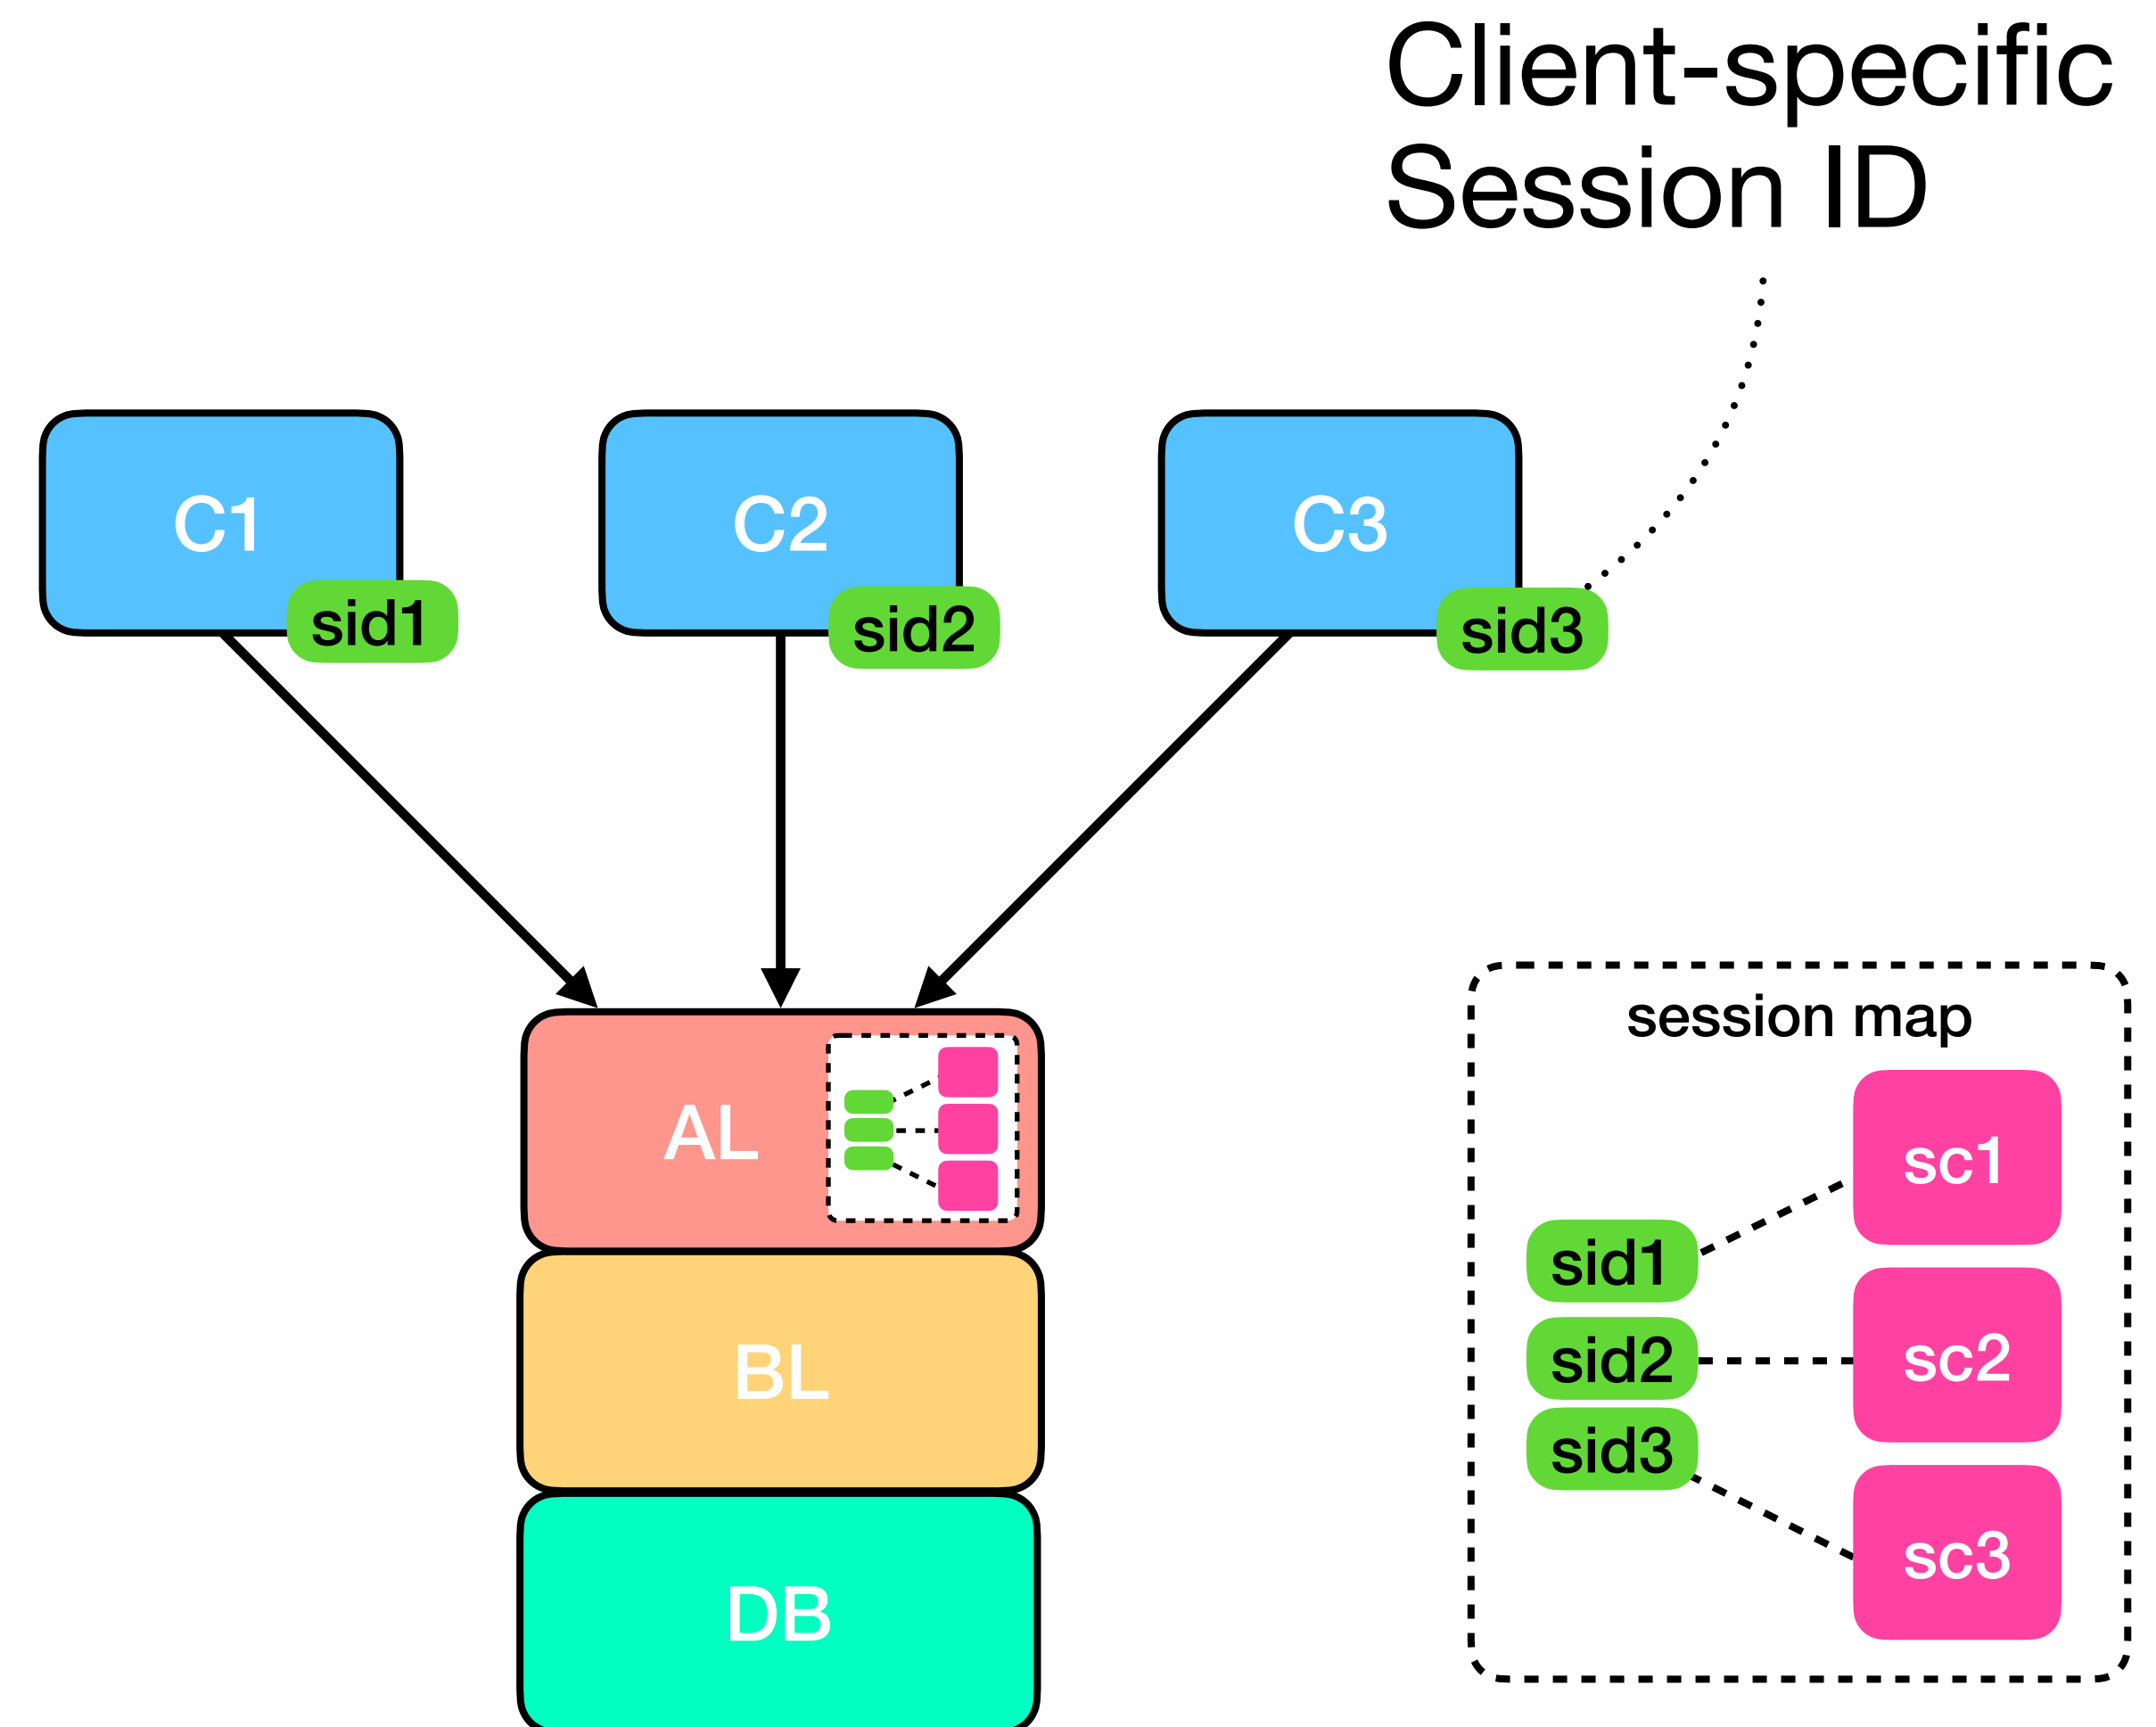
# Client-side session context

- Session context in **browser memory** (e.g., in cookies or local storage)
- **Advantages**
  - no memory consumption on server side, even if client breaks connection
  - server can be changed during session (→ load balancing later in this lesson)
  - session hijacking attacks can be harder
- **Disadvantages**
  - higher data transfer volume - complete context has to be transmitted
  - context may not survive client restarts



# Server-side session context

- Clients only store small keys (session IDs)
- Server holds **session map**  $\text{sid} \rightarrow \text{context}$ , *usually* in the AL/Web Server component
- **Advantages**
  - very little extra information transferred
  - session can survive restarts of client
- **Disadvantages**
  - memory consumption on the server side
  - switching servers during session requires special attendance
  - server has to monitor / evict inactive sessions
  - leaking session IDs can lead to hijacking





# Realizing Sessions

- In most systems, **server-side storage** of the context is used
- Identification of simultaneous sessions realized via **session IDs**
- Transfer of the session IDs with **each HTTP request**
  - as request **parameter** in **GET** or **POST** requests
  - as **part of the URL** (requires URL rewriting)
  - as **session cookie** in HTTP headers
- The session context **binds server resources** during the lifetime of a session
- Servers usually define a **session timeout** to be able to remove inactive sessions to reclaim and reuse those resources



# URL rewriting

- URL rewriting transmits all session-dependent data as **parameters** in the URL (XYZ should be random and hard to guess):

```
1 http://host/application/page.ext?SessionID=XYZ  
2 or  
3 http://host/application/XYZ/page.ext
```

- Requires **inspection** and **rewriting** of URLs in links inside **page content**

# Drawbacks of URL Rewriting

- the URL can become **messy** and **error-prone**
- **bookmarking** becomes impossible
- the approach can be unusable due to **length-limits** of URLs on some systems
- the URLs inside the documents have to be **adapted dynamically** to include the session ID
- it is **inherently insecure** (unless using https, or unless no content at web-space should ever be protected).

# Cookies

- Cookies are **small text files** used to store server information (e.g., a session ID) on the client node as **name-value pairs**.
- Cookies allow to make the session context at the server's side accessible by an ID via the session map.

# Cookies

- Web server **transmits cookies** to the browser in the HTTP **response header**
  - Browser **re-transmits cookies** to the respective server with each subsequent request in the **request header**
  - **Session cookies**
    - kept in browser memory
    - deleted as soon as the browser terminates
  - **Permanent cookies**
    - stored persistently on disk
    - for a certain lifetime or unlimited
- Cookies represent **sensitive information**
    - **identification of a user** during a session
    - countermeasures against **session hijacking attacks** have to be taken
    - **tracking** of activities possible
    - ...

# Cookie Headers

```
1  ----> GET ...
2
3  <---- HTTP/1.1 200 OK
4      Content-type: text/html
5      Set-Cookie: name=value
6      (content of page)
7
8  ----> GET ... HTTP/1.1
9      Host: www.uni-koblenz.de
10     Cookie: name=value
11     Accept: */*
```

# Examples of session cookies (using browser tools)

# 4.4 Vertical and Horizontal Scaling, Load Balancing



# Scalability...

- **Scalability** is a quality measure for **how easy a system can be adapted to varying loads**
- Adaption means to **react** on **higher** loads by **scaling up**, and to react on **lower** loads by **scaling down**
- Ideally, a system can be **scaled without interruption** of service
- In software systems, **architectural design** plays a **central role to achieve scalability**
- Dealing with web applications, we assume **distributed systems** consisting of web browsers, web servers, application servers, database servers, (any possibly many more components)
- Systems with **layered architectures** generally are **more flexible** since scaling can be applied on the bottlenecks only


# Resource Congestion

- Network and server resources **shared** by many clients
- Limitations due to
  - Network **bandwidth**
  - Number of **concurrent connections**
  - **Memory** consumption
  - **CPU** consumption
  - **Storage space**
  - **Storage bandwidth**
  - ...and other factors
- Increased number of concurrent clients and/or requests consume resources
- High load results in slow processing, or even in denial of service

# Vertical Scaling

Put more power to a machine



- Usable in situations of
  - Memory congestion
  - CPU congestion
  - Storage space congestion
  - Storage bandwidth congestion
- Not applicable to deal with limits in network bandwidth or concurrent connections



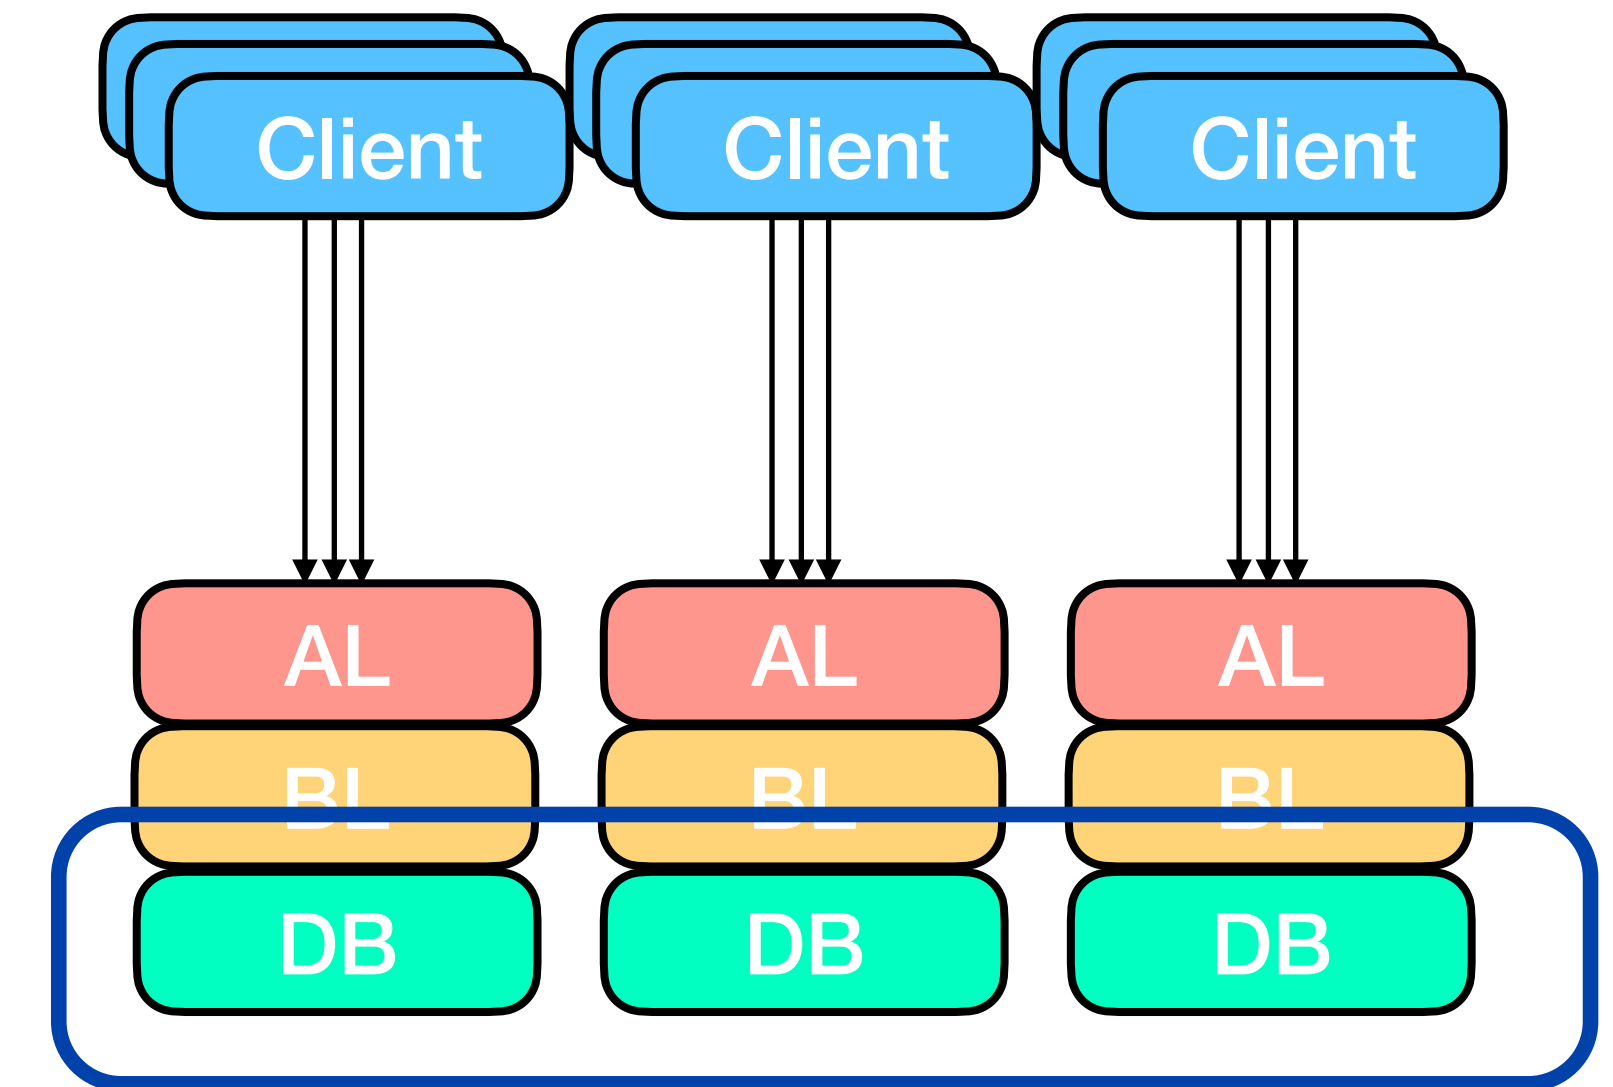
# Horizontal Scaling

Increase number of machines

- Usable in situations of
  - Network bandwidth congestion
  - Exceeded number of concurrent connections
  - Memory congestion
  - CPU congestion
  - Storage space congestion
  - Storage bandwidth congestion

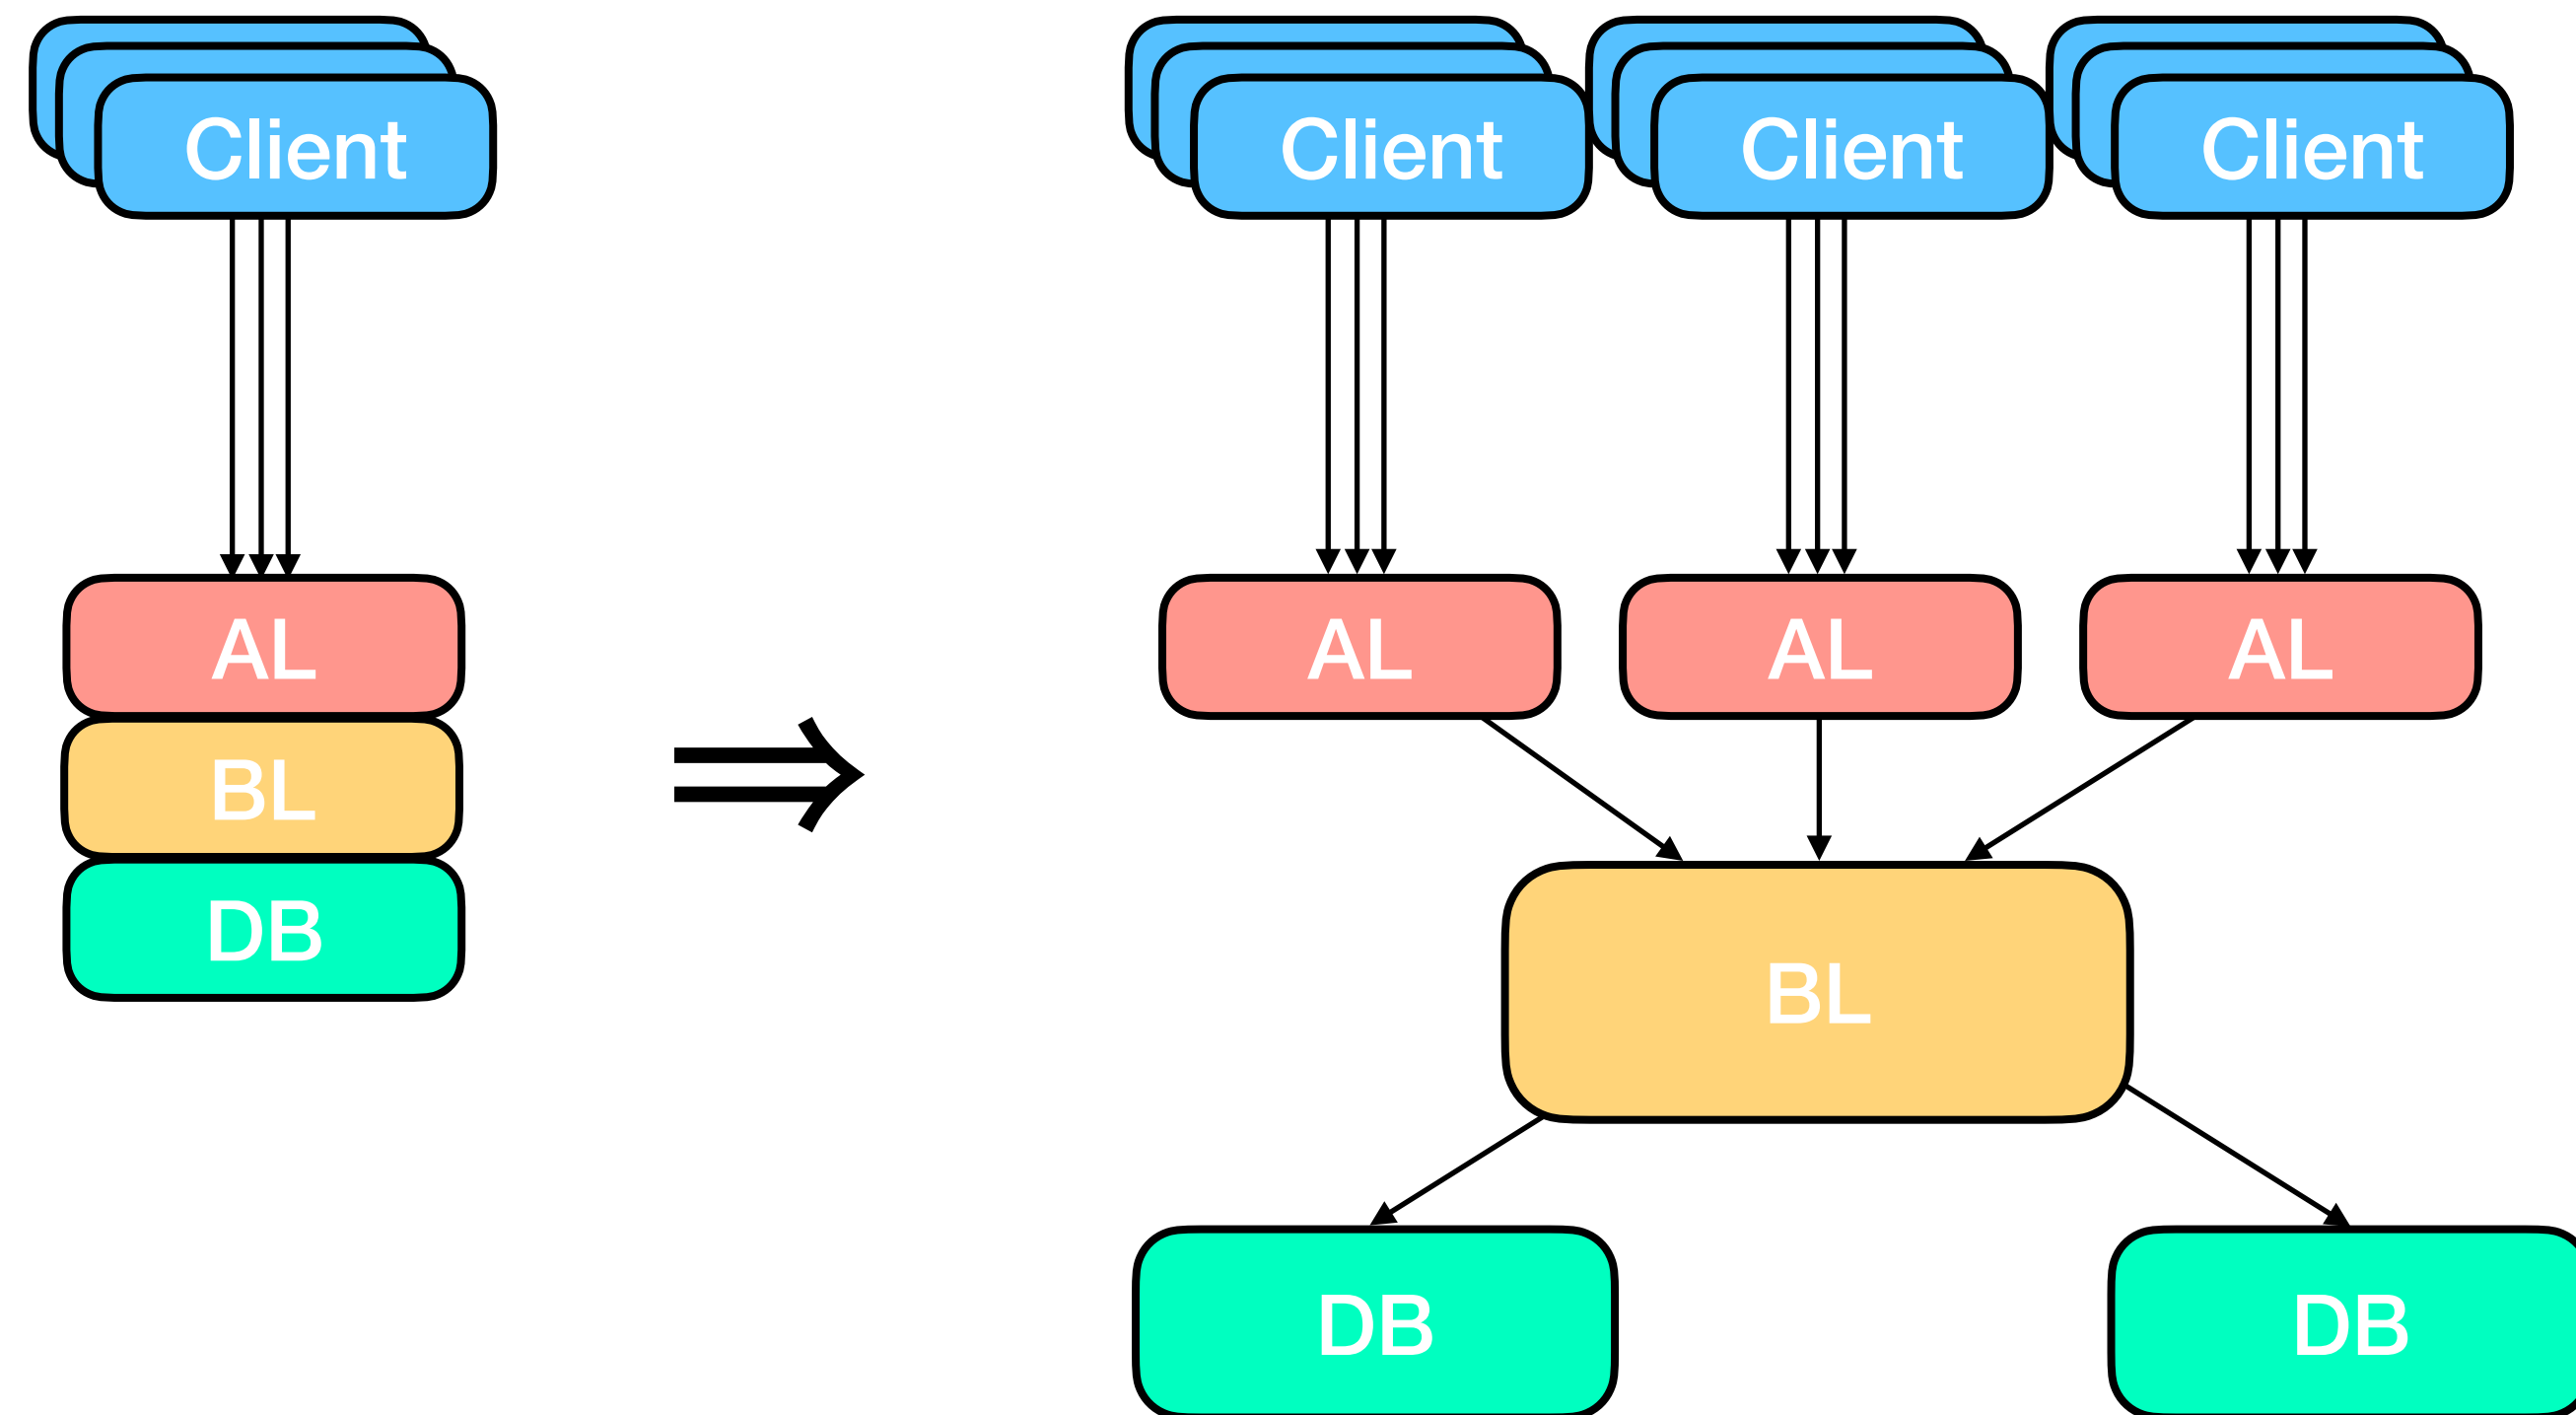


Load balancing required to distribute clients/requests to multiple servers



Horizontal DB scaling requires **separation** of and **synchronization** between data partitions

# Combined Vertical and Horizontal Scaling

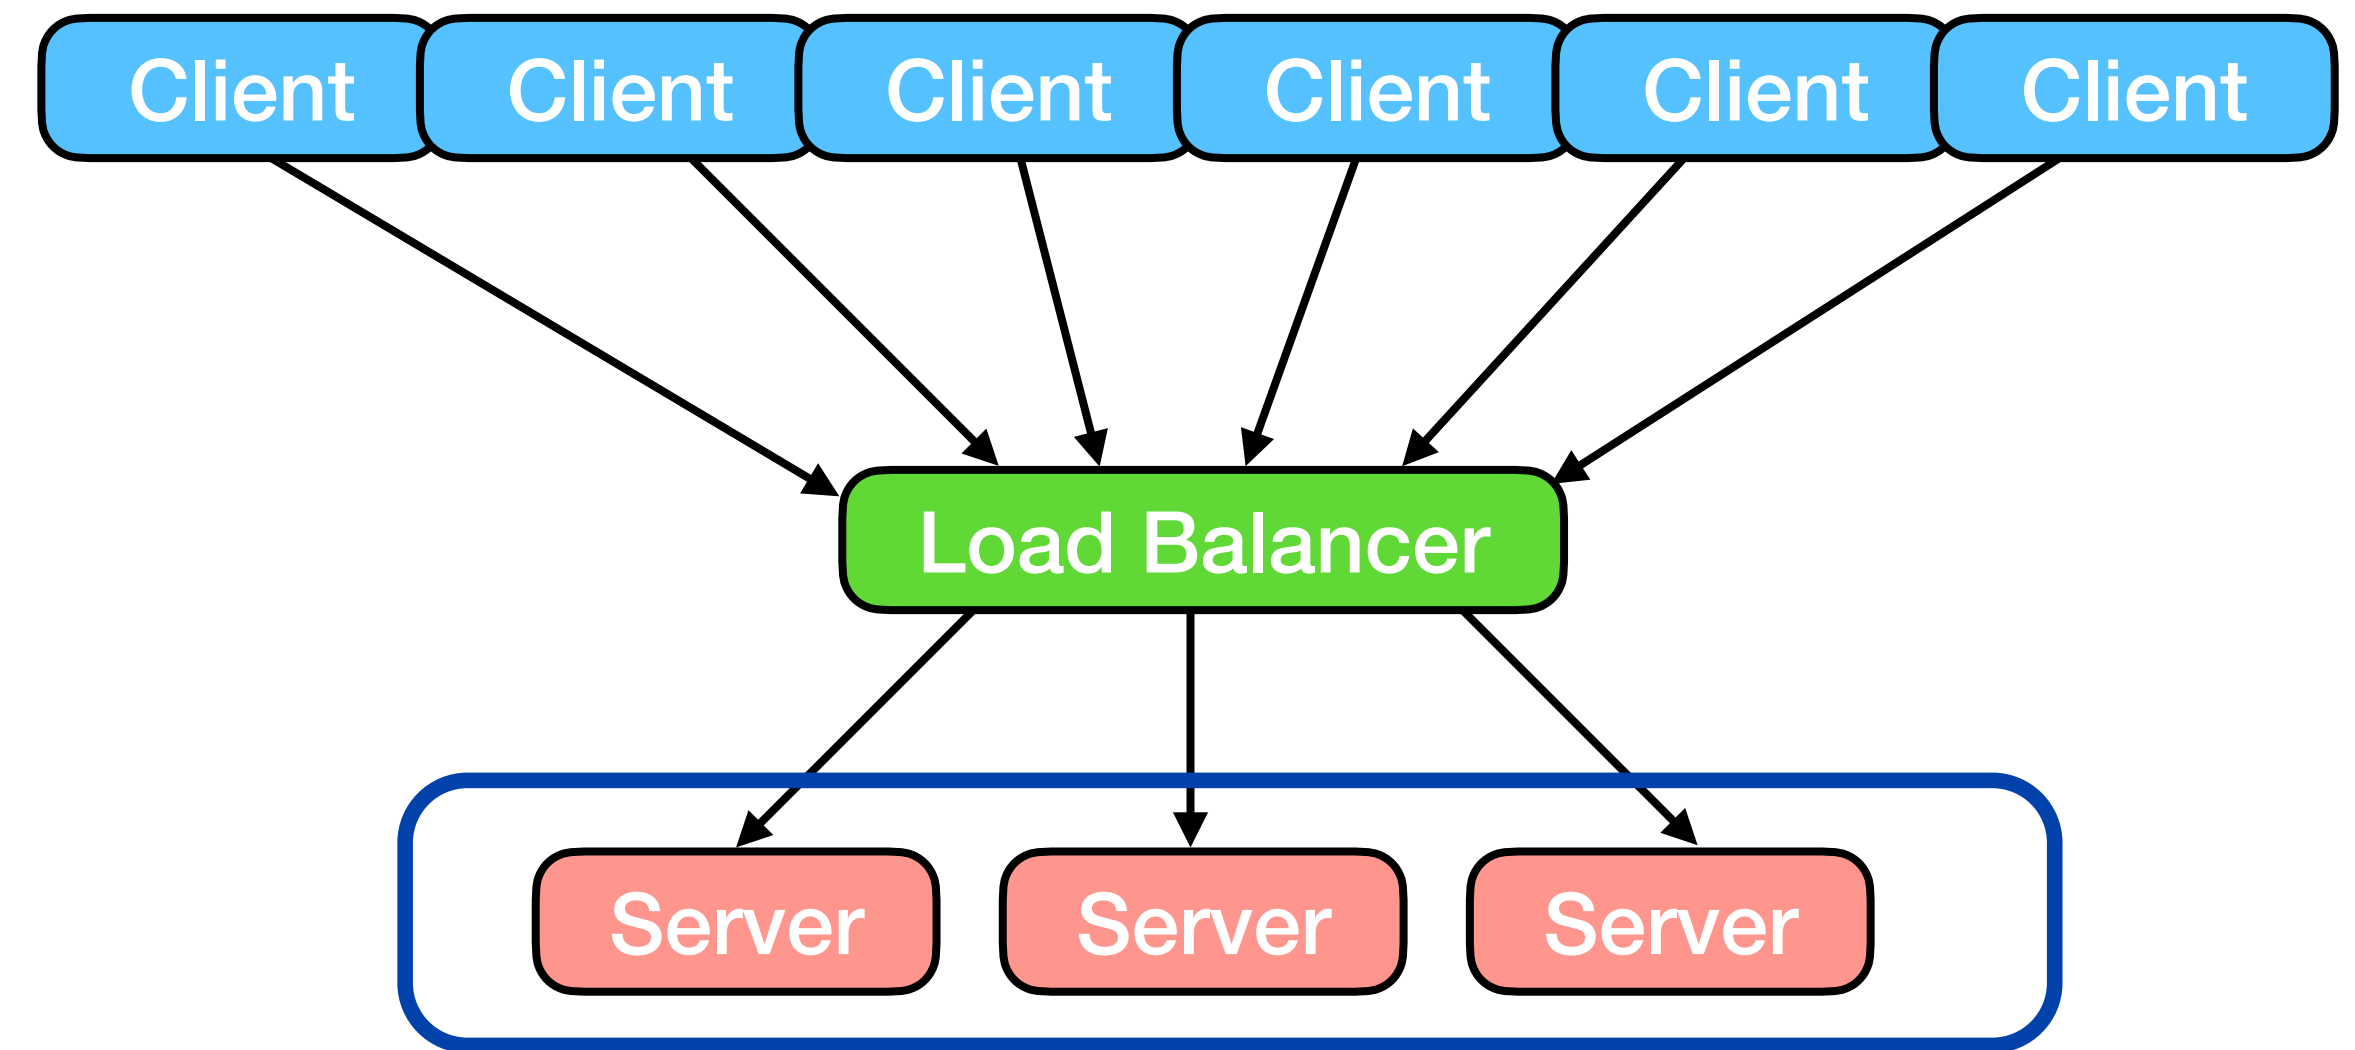


Layered architectures allow to **scale each layer individually** depending on the individual congestion problem



# Load Balancing

- Transparent to the clients - only a single server from a client's perspective
- Load balancing can work on various layers of the network stack (see page #7)
- Static and dynamic variants
  - Dynamic variants consider system state to take decision, static variant doesn't
- Methods include
  - DNS round robin (simplest approach)
  - NAT based load balancing (local, network layer)
  - Anycast load balancing (global, network layer)
  - Flat based balancing (local, link layer)
  - Application layer balancing



- Load balancer can become new bottleneck
- Data inspection can be infeasible due to end-to-end encryption

# Load Balancing

- DNS round robin
  - DNS = Domain Name System
  - Multiple servers with a single name
  - Name-to-IP-Address resolution results in different IP on each request
  - Easy implementation, on on-premise devices
  - Difficult to realize server-based session stores
  - Client “sees” multiple servers
- NAT based load balancing
  - NAT = Network Address Translation
  - Multiple servers on internal network with different addresses
  - Load balancer/router rewrites network packets from clients to reach different internal servers
  - All traffic needs to go through load balancer
  - Server-based session stores require “server affinity”, i.e., the same server has to be assigned to a client during a session
  - Servers appear as single node clients

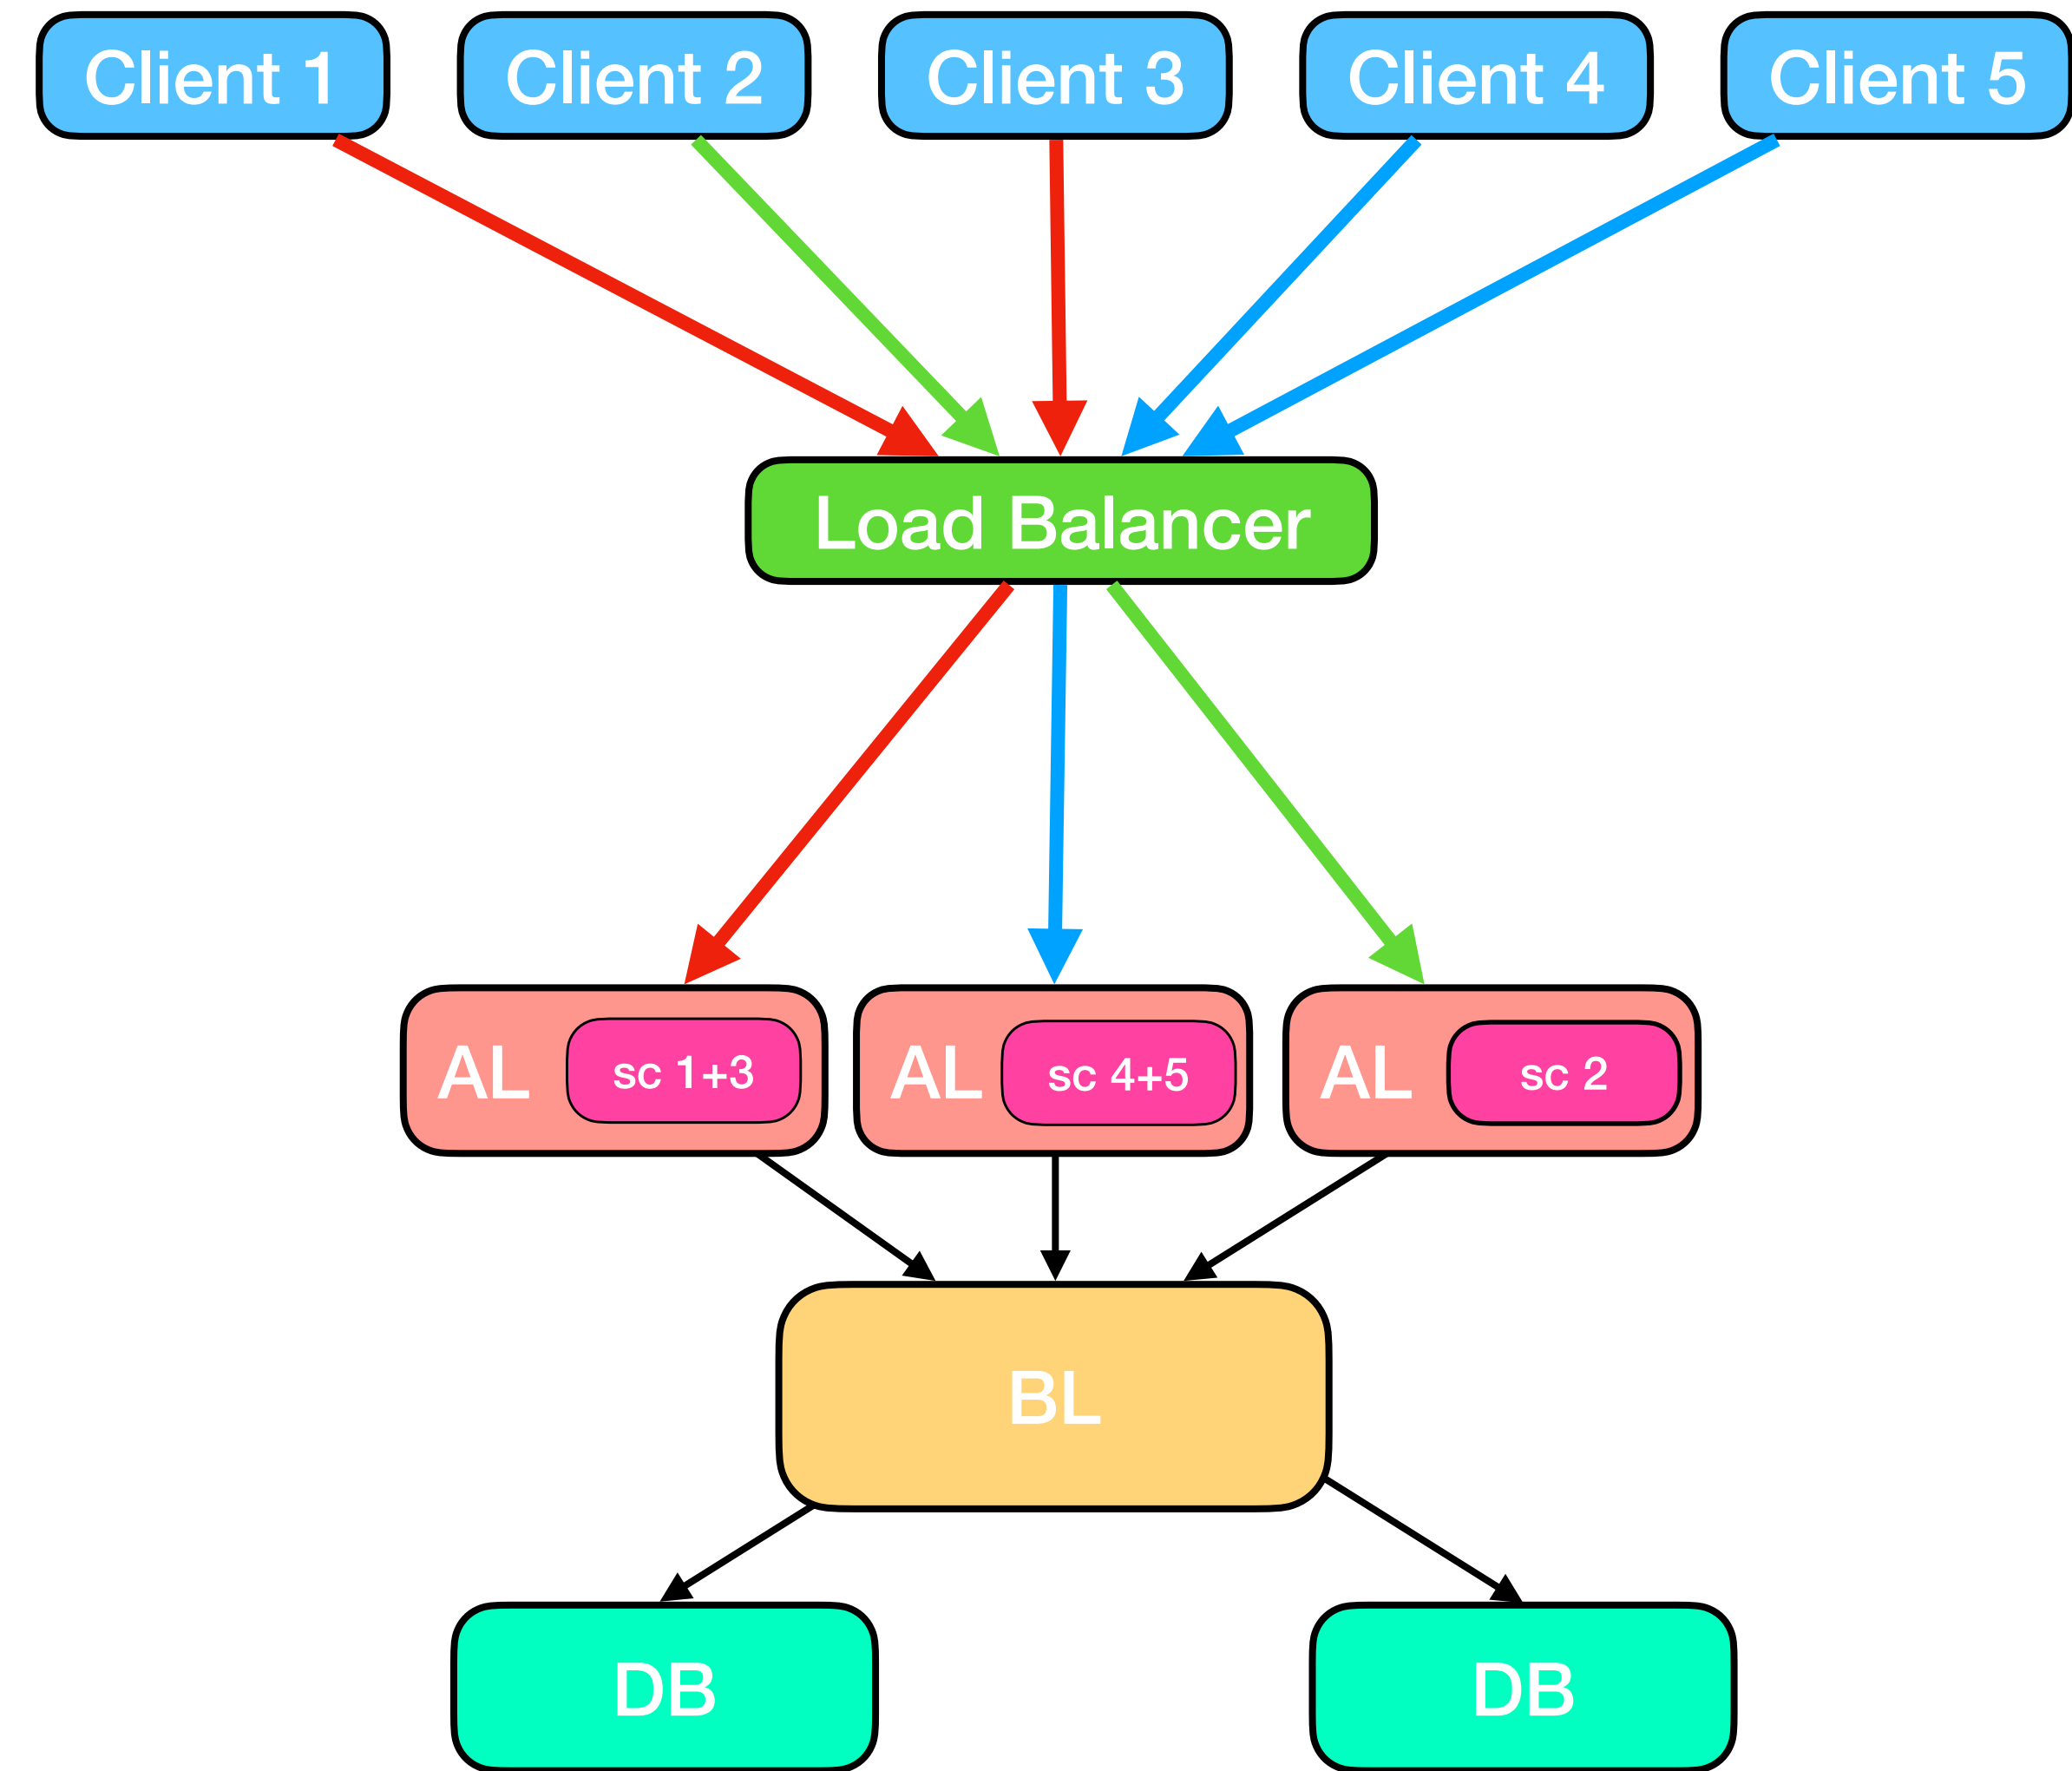


# Load Balancing

- Anycast load balancing
  - Multiple servers with same IP address in spatially separated networks
  - Global routing strategy (BGP, Border Gateway Protocol) ensures that only one server responds
  - Server with shortest network route (number of “hops”) is selected
  - Usually, picks a server geographically close to the client
  - Server-based sessions require relatively “stable” network
  - Servers appear as single node to clients
- Flat based balancing
  - Multiple servers with different addresses on local network
  - Also known as “MAT” - MAC Address Translation (MAC = Media Access Control address, physical local network address of a node)
  - Load balancer rewrites incoming packages to different server MACs, outgoing traffic from server to client bypasses load balancer
  - Beneficial because in general, incoming data volume is way smaller than outgoing; prevents overload of the load balancer
  - Affinity required for server-based sessions
  - Servers appear as single node to clients

# Load Balancing vs. Sessions

- Presence of server-side session context requires “**server affinity**”, i.e. the **same server** has to be used for **subsequent requests** of a client
- Depending on the scaling strategy, session maps have to be **moved to BL layer** and/or stored in the **DB layer**
  - Increased memory load to lower layers, results in slower session access
  - More traffic between lower layers
  - Possibly need to synchronize distributed databases and/or implement “affinity” on lower layers



# Stateless Interface

- Horizontal scaling is easier with **stateless** interfaces
- Stateless means that a server instance doesn't store any state information
- Hence, each server instance is equivalent to all others
- → Facade pattern in the architecture chapter
- Load balancer may select any server for each request, no affinity required
- Stateless AL and BL layers can be achieved by client-side session context, or by storing the session context in the DB layer
- As mentioned before, data transfer volume increases while memory load in AL and/or BL decreases
- Careful investigation of performance bottlenecks required to achieve an appropriate solution design in such tradeoff situations

# What we have learned...

## Communication (Part I)

- ✓ Network basics
- ✓ HTTP
- ✓ Sessions
- ✓ Scaling and Load Balancing




Image: colourbox.de